

## hw2\_liuke

October 28, 2018

```
In [63]: import numpy
        from urllib.request import urlopen
        import scipy.optimize
        import random
        from math import exp
        from math import log

In [64]: def parseData(fname):
          for l in urlopen(fname):
              yield eval(l)

In [65]: print("Reading data...")
          data = list(parseData("http://jmcauley.ucsd.edu/cse190/data/beer/beer_50000.json"))
          print("done")

Reading data...
done

In [66]: def feature(datum):
          feat = [1, datum['review/taste'], datum['review/appearance'], \
                  datum['review/aroma'], datum['review/palate'], datum['review/overall']]
          return feat

In [67]: def inner(x,y):
          return sum([x[i]*y[i] for i in range(len(x))])

          def sigmoid(x):
              return 1.0 / (1 + exp(-x))

In [68]: ##### # Logistic regression by gradient ascent      #
##### # NEGATIVE Log-likelihood
def f(theta, X, y, lam):
    loglikelihood = 0
    for i in range(len(X)):
        logit = inner(X[i], theta)
```

```

loglikelihood -= log(1 + exp(-logit))
if not y[i]:
    loglikelihood -= logit
for k in range(len(theta)):
    loglikelihood -= lam * theta[k]*theta[k]
# for debugging
# print("ll =" + str(loglikelihood))
return -loglikelihood

# NEGATIVE Derivative of log-likelihood
def fprime(theta, X, y, lam):
    dl = [0]*len(theta)
    for i in range(len(X)):
        logit = inner(X[i], theta)
        for k in range(len(theta)):
            dl[k] += X[i][k] * (1 - sigmoid(logit))
            if not y[i]:
                dl[k] -= X[i][k]
    for k in range(len(theta)):
        dl[k] -= lam*2*theta[k]
    return numpy.array([-x for x in dl])

```

## 0.1 Problem 1

```

In [69]: import random
         random.shuffle(data)
         X = [feature(d) for d in data]
         y = [d['beer/ABV'] >= 6.5 for d in data]

In [70]: X_train = X[:len(X)//3]      # top 1/3 of the shuffled data
         y_train = y[:len(y)//3]
         #####
         # Train
         #####
         def train(lam):
             theta,_,_ = scipy.optimize.fmin_l_bfgs_b(f, [0]*len(X[0]), fprime, pgtol = 10, args
             return theta

In [71]: ##### # Predict #####
         #####
         def performance(theta,X,Y):
             scores = [inner(theta,x) for x in X]
             predictions = [s > 0 for s in scores]
             correct = [(a == b) for (a,b) in zip(predictions,Y)]
             acc = sum(correct) * 1.0 / len(correct)
             return acc

```

```
In [72]: #####
# Validation pipeline
#####
### training set
lam = 1.0

theta = train(lam)
acc = performance(theta,X_train,y_train)
print("lambda = " + str(lam) + ":\taccuracy of the training set=" + str(acc))

lambda = 1.0:      accuracy of the training set=0.7224888995559823
```

In [73]: ### validation set

```
X_validation=X[len(X)//3:2*len(X)//3]
y_validation=y[len(y)//3:2*len(y)//3]
acc1 = performance(theta,X_validation,y_validation)
print("lambda = " + str(lam) + ":\taccuracy of the validation set=" + str(acc1))

lambda = 1.0:      accuracy of the validation set=0.7171656566868663
```

In [74]: ### test set

```
X_test=X[2*len(X)//3:]
y_test=y[2*len(y)//3:]
acc2 = performance(theta,X_test,y_test)
print("lambda = " + str(lam) + ":\taccuracy of the validation set=" + str(acc2))

lambda = 1.0:      accuracy of the validation set=0.7178856422871542
```

## 0.2 Problem 2

```
In [75]: ### P/N/TP/TN/FP/FN of the test set
scores2 = [inner(theta,x) for x in X_test]
predictions2 = [s > 0 for s in scores2]

### positives
nums_positive = sum(predictions2)
print("The number of predicted positives:\t" + str(nums_positive))

### negatives
nums_negative = len(predictions2)-nums_positive
print("The number of predicted negatives:\t" + str(nums_negative))

### True Positives
TP = [(a == b and a == True) for (a,b) in zip(predictions2,y_test)]
```

```

nums_TP = sum(TP)
print("The number of True Positives:\t" + str(nums_TP))

### True Negatives
TN = [(a == b and a == False) for (a,b) in zip(predictions2,y_test)]
nums_TN = sum(TN)
print("The number of True Negatives:\t" + str(nums_TN))

### False Positives
FP = [(a == True and b == False) for (a,b) in zip(predictions2,y_test)]
nums_FP = sum(FP)
print("The number of False Positives:\t" + str(nums_FP))

### False Negatives
FN = [(a == False and b == True) for (a,b) in zip(predictions2,y_test)]
nums_FN = sum(FN)
print("The number of False Negatives:\t" + str(nums_FN))

The number of predicted positives:      12350
The number of predicted negatives:      4317
The number of True Positives:          9038
The number of True Negatives:          2927
The number of False Positives:         3312
The number of False Negatives:         1390

```

### 0.3 Problem 3

If we want to assign greater importance to False Positives as compared to False Negatives, we could reduce the likelihood of FP, that is, case  $y>0$  and  $\text{logit}<0$ .

```

In [86]: #####
# Logistic regression by gradient ascent      #
######
# NEGATIVE Log-likelihood
def f_weighted(theta, X, y, lam):
    loglikelihood = 0
    for i in range(len(X)):
        logit = inner(X[i], theta)
        loglikelihood -= log(1 + exp(-logit))
        if not y[i]:
            if logit>0 :
                loglikelihood -= logit*1.4 ## modified place
                # loglikelihood *= 10
            else :
                loglikelihood -= logit
    for k in range(len(theta)):
        loglikelihood -= lam * theta[k]*theta[k]

```

```

# for debugging
# print("ll =" + str(loglikelihood))
return -loglikelihood

# NEGATIVE Derivative of log-likelihood
def fprime_weighted(theta, X, y, lam):
    dl = [0]*len(theta)
    for i in range(len(X)):
        logit = inner(X[i], theta)
        for k in range(len(theta)):
            dl[k] += X[i][k] * (1 - sigmoid(logit))
        if not y[i]:
            if logit >0 :
                dl[k] -= X[i][k]*1.4
                # dl[k] *= 10
            else :
                dl[k] -= X[i][k]
    for k in range(len(theta)):
        dl[k] -= lam*2*theta[k]
    return numpy.array([-x for x in dl])

def train_weighted(lam):
    theta,_,_ = scipy.optimize.fmin_l_bfgs_b(f_weighted, [0]*len(X[0]), fprime_weighted
                                              args = (X_train, y_train, lam))
    return theta

```

```

In [87]: theta1 = train_weighted(lam)
scores_weighted = [inner(theta1,x) for x in X_test]
predictions_weighted = [s > 0 for s in scores_weighted]

### True Positives
TP1 = [(a == b and a == True) for (a,b) in zip(predictions_weighted,y_test)]
nums_TP1 = sum(TP1)
print("The number of True Positives:\t" + str(nums_TP1))

### True Negatives
TN1 = [(a == b and a == False) for (a,b) in zip(predictions_weighted,y_test)]
nums_TN1 = sum(TN1)
print("The number of True Negatives:\t" + str(nums_TN1))

### False Positives
FP1 = [(a == True and b == False) for (a,b) in zip(predictions_weighted,y_test)]
nums_FP1 = sum(FP1)
print("The number of False Positives:\t" + str(nums_FP1))

### False Negatives
FN1 = [(a == False and b == True) for (a,b) in zip(predictions_weighted,y_test)]
nums_FN1 = sum(FN1)

```

```

print("The number of False Negatives:\t" + str(nums_FN1))

The number of True Positives:      8418
The number of True Negatives:      3634
The number of False Positives:     2605
The number of False Negatives:     2010

```

## 0.4 Problem 4

```

In [88]: ##### lam = 0
          theta_0 = train(0)
          acc_0 = performance(theta_0,X_validation,y_validation)
          print("lambda = " + str(0) + ":\taccuracy of the validation set=" + str(acc_0))

          ### lam = 0.01
          theta_1 = train(0.01)
          acc_1 = performance(theta_1,X_validation,y_validation)
          print("lambda = " + str(0.01) + ":\taccuracy of the validation set=" + str(acc_1))

          ### lam = 0.1
          theta_2 = train(0.1)
          acc_2 = performance(theta_2,X_validation,y_validation)
          print("lambda = " + str(0.1) + ":\taccuracy of the validation set=" + str(acc_2))

          ### lam = 1
          theta_3 = train(1)
          acc_3 = performance(theta_3,X_validation,y_validation)
          print("lambda = " + str(1) + ":\taccuracy of the validation set=" + str(acc_3))

          ### lam = 100
          theta_4 = train(100)
          acc_4 = performance(theta_4,X_validation,y_validation)
          print("lambda = " + str(100) + ":\taccuracy of the validation set=" + str(acc_4))

lambda = 0:      accuracy of the validation set=0.7177056458870823
lambda = 0.01:    accuracy of the validation set=0.7177056458870823
lambda = 0.1:     accuracy of the validation set=0.7177656446871062
lambda = 1:       accuracy of the validation set=0.7171656566868663
lambda = 100:     accuracy of the validation set=0.6693466130677387

```

The best  $\lambda$  is 100

```

In [89]: acc_train = performance(theta_4,X_train,y_train)
          print("lambda = " + str(100) + ":\taccuracy of the train set=" + str(acc_train))
          acc_test = performance(theta_4,X_test,y_train)
          print("lambda = " + str(100) + ":\taccuracy of the test set=" + str(acc_test))

```

```
lambda = 100:           accuracy of the train set=0.6740669626785072
lambda = 100:           accuracy of the test set=0.5965438617544702
```

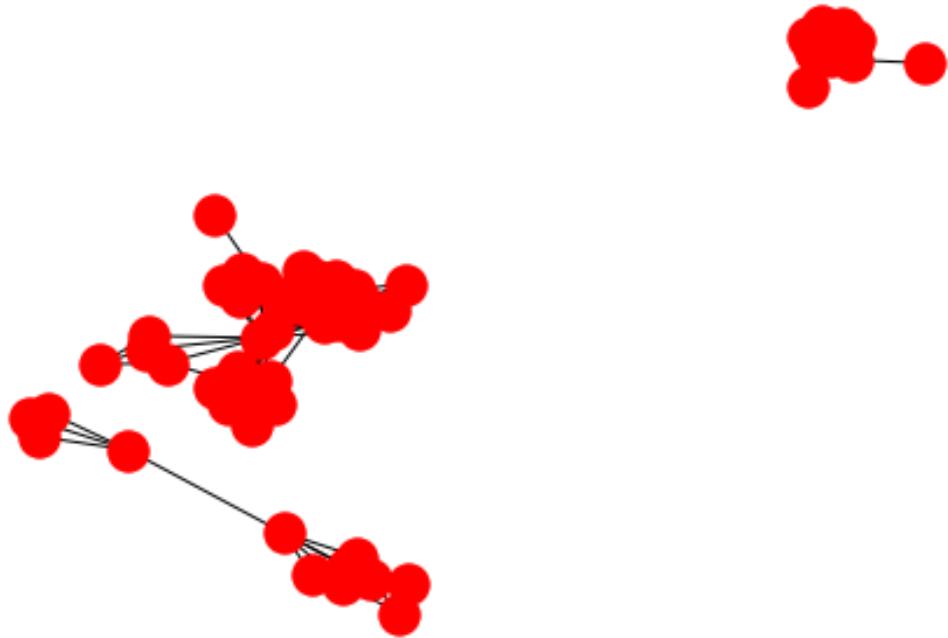
## 1 Task 2

### 1.1 problem 5

```
In [40]: ### Network visualization ###
    import networkx as nx
    import matplotlib.pyplot as plt
```

```
In [41]: edges = set()
nodes = set()
for edge in urlopen("http://jmcauley.ucsd.edu/cse258/data/facebook/egonet.txt"):
    x,y = edge.split()
    x,y = int(x),int(y)
    edges.add((x,y))
    edges.add((y,x))
    nodes.add(x)
    nodes.add(y)
```

```
In [42]: G = nx.Graph()
for e in edges:
    G.add_edge(e[0],e[1])
nx.draw(G)
plt.show()
plt.clf()
```



<Figure size 432x288 with 0 Axes>

```
In [43]: ### define a dict :group  [ key: node i; value: all node j if (i,j) in edges]
group={}
for edge in edges :
    if edge[0] in group :
        group[edge[0]].add(edge[1])
    else :
        group[edge[0]]=set([edge[1]])

### use BFS to find all nodes that belongs to one group
def get_all_connected_groups(graph):
    already_seen = set()
    result = []
    for node in graph:
        if node not in already_seen:
            connected_group, already_seen = \
            get_connected_group(graph,node, already_seen)
            result.append(connected_group)
    return result
```

```

def get_connected_group(graph,node, already_seen):
    result = []
    nodes = set([node])
    while nodes:
        node = nodes.pop()
        already_seen.add(node)
        nodes.update(graph[node] - already_seen)
        result.append(node)
    return result, already_seen

components = get_all_connected_groups(group)
#print(components)

print("\n The number of connected components:\t" +str(len(components)))
length = [0]*len(components)
for i in range(len(components)):
    length[i] = len(components[i])
print("The largest component's node number\t"+str(max(length)))

```

The number of connected components:	3
The largest component's node number	40

## 1.2 Problem 6

In [44]: *### new\_graph is the largest connected component*

```

new_graph = components[0]
new_graph.sort()

### split new_graph into two parts
graph1 = new_graph[:len(new_graph)//2]
graph2 = new_graph[len(new_graph)//2:]

### find the normalized cut
def norm_cut(graph1,graph2) :
    cut = 0
    degree1 = sum(G.degree(node1) for node1 in graph1)
    degree2 = sum(G.degree(node2) for node2 in graph2)
    for node1 in graph1 :
        for node2 in graph2 :
            if (node1,node2) in edges:
                cut+=1
    Norm_cut=0.5 * (cut/degree1 + cut/degree2)
    return Norm_cut
res = norm_cut(graph1,graph2)
print("The normalized cut cost is: "+ str(res))

```

The normalized cut cost is: 0.42240587695133147

### 1.3 Problem 7

```
In [50]: # minimize the cost, if new generated cost is higher than the
# previous one, then stop

old_min_cost = res
new_min_cost = res
while True :
    if old_min_cost < new_min_cost :
        break
    else :
        old_min_cost = new_min_cost

cost = [0] * len(new_graph)
for i in range(len(new_graph)) :
    g1= graph1.copy()
    g2= graph2.copy()
    if new_graph[i] in g1 :
        g2 += [new_graph[i]]
        del g1[g1.index(new_graph[i])]
    else :
        g1 += [new_graph[i]]
        del g2[g2.index(new_graph[i])]
    cost[i] = norm_cut (g1, g2)

move_point = cost.index(min(cost))

if new_graph[move_point] in graph1 :
    graph2 += [new_graph[move_point]]
    del graph1[graph1.index(new_graph[move_point])]
else :
    graph1 += [new_graph[move_point]]
    del graph2[graph2.index(new_graph[move_point])]
new_min_cost = min(cost)

print("The normalized cut cost is "+str(old_min_cost))
print("nodes in Community 1: "+str(graph1))
print("nodes in Community 2: "+str(graph2))

### because the content is cut when saving as PDF,
### I will show the result in annotation
### community 1:[697, 703, 708, 713, 719, 745, 747, 753, 769, 772, 774,
### 800, 803, 805, 810, 811, 819, 828, 823, 830, 840, 880, 890, 869, 856]
### community 2: [825, 861, 863, 864, 876, 878, 882, 884,
### 886, 888, 889, 893, 729, 804, 798]
```

```
The normalized cut cost is 0.09817045961624274
```

```
nodes in Community 1: [697, 703, 708, 713, 719, 745, 747, 753, 769, 772, 774, 800, 803, 805, 8]
```

```
nodes in Community 2: [825, 861, 863, 864, 876, 878, 882, 884, 886, 888, 889, 893, 729, 804, 79]
```

## 1.4 Problem 8

```
In [55]: ## new_graph : 40 nodes : a connected tcommunity
## all_edge : number of edges in the new_graph
all_edge = sum([G.degree(v) for v in new_graph])/2
#print(all_edge)

## split into two parts
graph_1 = new_graph[:len(new_graph)//2]
graph_2 = new_graph[len(new_graph)//2:]

### compute Q
def q(all_edge, g1, g2):
    count = 0
    for i in g1 :
        for j in g2 :
            if (i,j) in edges :
                count += 1
    q_=(all_edge-count)/all_edge
    a = sum([G.degree(v) for v in g1])
    q_ -= (a/(all_edge*2))**2
    a = sum([G.degree(v) for v in g2])
    q_ -= (a/(all_edge*2))**2
    return q_

### Greedy algorithm , when new Q is not increasing, stop.
old_q = new_q = -0.5
while True :
    if old_q > new_q :
        break
    else :
        old_q = new_q

    q_list = [0] * len(new_graph)
    for i in range(len(new_graph)) :
        g1 = graph_1.copy()
        g2 = graph_2.copy()
        if new_graph[i] in g1 :
            g2 += [new_graph[i]]
            del g1[g1.index(new_graph[i])]
        else :
            g1 += [new_graph[i]]
            del g2[g2.index(new_graph[i])]
```

```

q_list[i] = q(all_edge, g1, g2)

mp= q_list.index(max(q_list))
if new_graph[mp] in graph_1 :
    graph_2 += [new_graph[mp]]
    del graph_1[graph_1.index(new_graph[mp])]
else :
    graph_1 += [new_graph[mp]]
    del graph_2[graph_2.index(new_graph[mp])]
new_q = max(q_list)
print("Community 1:"+str(graph1))
print("Community 2:"+str(graph2))
print("modularity values for the split: "+str(old_q))

### the two community we found here is the same with that in Problem 7

```

```

Community 1:[697, 703, 708, 713, 719, 745, 747, 753, 769, 772, 774, 800, 803, 805, 810, 811, 812, 813, 814, 815, 816, 817, 818, 819, 820, 821, 822, 823, 824, 825]
Community 2:[825, 861, 863, 864, 876, 878, 882, 884, 886, 888, 889, 893, 729, 804, 798]
modularity values for the split: 0.3380165289256196

```

In [ ]: